



**EPAM Systems, RD Dep., RD Dep.**

# **POSTGRESQL DB FOR DWH AND ETL BUILDING**

---

## **PostgreSQL Relational Structures**

**Legal Notice:** This document contains privileged and/or confidential information and may not be disclosed, distributed or reproduced without the prior written permission of EPAM®.

---

Confidential

CONTENTS

1. PREREQUISITE TASK.....3

2. TABLES .....3

2.1 Task 1 - PERFORMANCE WITH UNLOGGED TABLES.....3

2.2 Task 2 - UNDERSTANDING INHERITED TABLES .....3

3. INDEXES .....4

3.1 Task 3 - BRIN vs B-TREE .....4

3.2 Task 4 - GIN vs GiST.....5

4. FOREIGN DATA WRAPPERS .....6

4.1 Task 5 - CSV FILE AS A TABLE.....6

## 1. PREREQUISITE TASK

Connect to your test\_db and lab schema which was created on the previous module. Check search path parameter, add your schema into path if it is needed.

Please DROP a table person from previous module and recreate it. Fill the person table with data:

```
CREATE TABLE labs.person (  
    id integer NOT NULL,  
    name varchar(15)  
);  
  
INSERT INTO person VALUES (1, 'Bob');  
INSERT INTO person VALUES (2, 'Alice');  
INSERT INTO person VALUES (3, 'Robert');
```

Read about EXPLAIN ANALYZE command. Do not dive deep for now.

## 2. TABLES

### 2.1 TASK 1 – PERFORMANCE WITH UNLOGGED TABLES

Task Result: Add your personal description of what happened and why with screenshots where needed.

1. Create simple table

```
labs.test_simple(a int,b int)
```

2. Perform insert and check how much time it takes:

```
insert into test_simple values (generate_series(1,1000000));
```

then

```
insert into test_simple values (generate_series(1,5000000));
```

3. Create UNLOGGED table

```
test_unlogged(a int,b int)
```

4. Perform insert and check how much time it takes:

```
insert into test_unlogged values (generate_series(1,1000000));
```

then

```
insert into test_unlogged values (generate_series(1,5000000));
```

### 2.2 TASK 2 – UNDERSTANDING INHERITED TABLES

Task Result: Provide queries where needed. Add your personal description of what happened and why with screenshots where needed.

1. Create inherited table for person:

```
CREATE TABLE labs.users(user_id INT GENERATED BY DEFAULT AS IDENTITY  
PRIMARY KEY, login VARCHAR(30)) INHERITS (person);
```

2. Perform insert into new table:

```
INSERT INTO users VALUES (1, 'new Bob', 'NewLogin');  
INSERT INTO users VALUES (999, 'TestUser', 'TestLogin');
```

3. Perform selects:

```
SELECT * FROM users;  
SELECT * FROM person;  
SELECT * FROM ONLY person;
```

4. Modify data in person table, check result:

```
UPDATE person  
SET name = 'not Bob'  
where id = 1;
```

Rewrite UPDATE to change name for row where id=1 for person table only.

5. Alter person table and check structure and constraints on both users and person tables:

```
ALTER TABLE person ADD COLUMN status integer DEFAULT 0;  
ALTER TABLE person ADD CONSTRAINT status CHECK (status in (0,1)) NO  
INHERIT;  
ALTER TABLE person ADD CONSTRAINT id UNIQUE (id, name);
```

## 3. INDEXES

### 3.1 TASK 3 - BRIN VS B-TREE

Task Result: Provide queries where needed. Add your personal description of what happened and why with screenshots where needed\*.

\* You could also use EXPLAIN ANALYZE for SELECT queries and check the plans differences.

1. Create table test\_index:

```
CREATE TABLE labs.test_index (  
    num float NOT NULL,  
    load_date timestampz NOT NULL  
);
```

2. Fill the table with a lot of test data:

```
INSERT INTO test_index(num, load_date)
SELECT random(), x
FROM generate_series('2017-01-01 0:00'::timestampz,
    '2021-12-31 23:59:59'::timestampz, '10 seconds'::interval) x;
```

Check the table size:

```
SELECT pg_size_pretty(pg_relation_size('test_index'));
```

3. Perform select and check how much time it takes:

```
SELECT date_trunc('year', load_date), max(num)
FROM test_index
WHERE load_date BETWEEN '2021-09-01 0:00' AND '2021-10-31 11:59:59'
GROUP BY 1
ORDER BY 1;
```

4. Create B-Tree Index on test\_index table for load\_date column. How long is this operation take? Repeat step 3. Check the index size(see step 1, change 'test\_index' to name of your index). Drop the B-Tree index.
5. Create BRIN Index on test\_index table for load\_date column. How long is this operation take? Repeat step 3. Check the index size(see step 1, change 'test\_index' to name of your index). Drop the BRIN index.
6. DROP test\_index table.

## 3.2 TASK 4 - GIN VS GIST

Task Result: Provide queries where needed. Add your personal description of what happened and why with screenshots where needed.

1. Create a table and fill with test data:

```
CREATE TABLE test_index AS SELECT id, md5(id::text) as t_hash
FROM generate_series(1, 10000000) AS id;
```

Check the table size(see Task 3).

2. Perform select and check how much time it takes:

```
SELECT * FROM test_index WHERE t_hash LIKE '%ceea167a5a%';
```

3. Create GIST Index on test\_index table for t\_hash column. How long is this operation take? Repeat step 1. Check the index size(see Task 3). Drop the GIST index.

```
CREATE EXTENSION pg_trgm;
```

```
CREATE INDEX idx_text_index_gist ON test_index USING gist(t_hash
gist_trgm_ops);
```

4. Create GIN Index on test\_index table for t\_hash column. How long is this operation take? Repeat step 1. Check the index size(see Task 3). Drop the GIN index.

```
CREATE INDEX idx_text_index_gin ON test_index USING gin (t_hash  
gin_trgm_ops);
```

5. DROP test\_index table.

## 4. FOREIGN DATA WRAPPERS

### 4.1 TASK 5 – CSV FILE AS A TABLE

Task Result: Provide queries where needed. Add your personal description of what happened and why with screenshots where needed.

1. Install file\_fdw extension to database and create the SERVER to use, every FOREIGN TABLE requires a server:

```
CREATE EXTENSION file_fdw;  
CREATE SERVER test_import FOREIGN DATA WRAPPER file_fdw;
```

2. Create the foreign table to connect to the CSV file (cities\_list.csv):

```
CREATE FOREIGN TABLE labs.test_foreign_table  
(  
    LatD INT,  
    LatM INT,  
    LatS INT,  
    NS TEXT,  
    LonD INT,  
    LonM INT,  
    LonS INT,  
    EW TEXT,  
    City TEXT,  
    State TEXT  
)  
SERVER test_import  
OPTIONS  
(  
    filename '\path_to_file\cities_list.csv',  
    format 'csv',  
    header 'true',  
    delimiter ','  
);
```

3. Select data from the foreign table and count of rows.

```
SELECT * FROM test_foreign_table;  
SELECT count(*) FROM test_foreign_table;
```

4. Create Materialized view to select from this foreign table. Select from MView.

```
select count(*) from mview
```

5. Delete from file some rows. Select count again.

6. Refresh MView to update count of rows.